



15 edycja konferencji SQLDay

8-10 maja 2023, WROCŁAW + ONLINE



partner złoty

Future Processing

partner srebrny



partner brązowy





Kamil Nowinski



Delta Lake Tables 101

10.05.2023

SQLDay 2023



Kamil Nowiński



Microsoft Data Platform **MVP**
Speaker, blogger, data enthusiast

Group Manager at Avanade UK&I (www.avanade.com)

>20 yrs experience as DEV/BI/(DBA)

Member of the Data Community PL

Founder of blog SQLPlayer (www.AzurePlayer.net)

GitHub: #adftools, SCD Merge Wizard and more...

SQL Server Certificates:

MCITP, MCP, MCTS, MCSA, MCSE Data Platform,

MCSE Data Management & Analytics, DevOps Expert

Moreover: Bicycle, Running, Digital photography

@NowinskiK, @Azure_Player



SQLPlayer is renaming



Blog

- Technical posts
- Various skill level
- Cheat sheets
- Recommended books
- Many useful other links
- Interviews (Podcast)
- YouTube Channel:
www.AzurePlayer.net/YouTube



Azure Player
Play with data & have fun!

www.AzurePlayer.net





Slides available:



<https://azureplayer.net/slides>

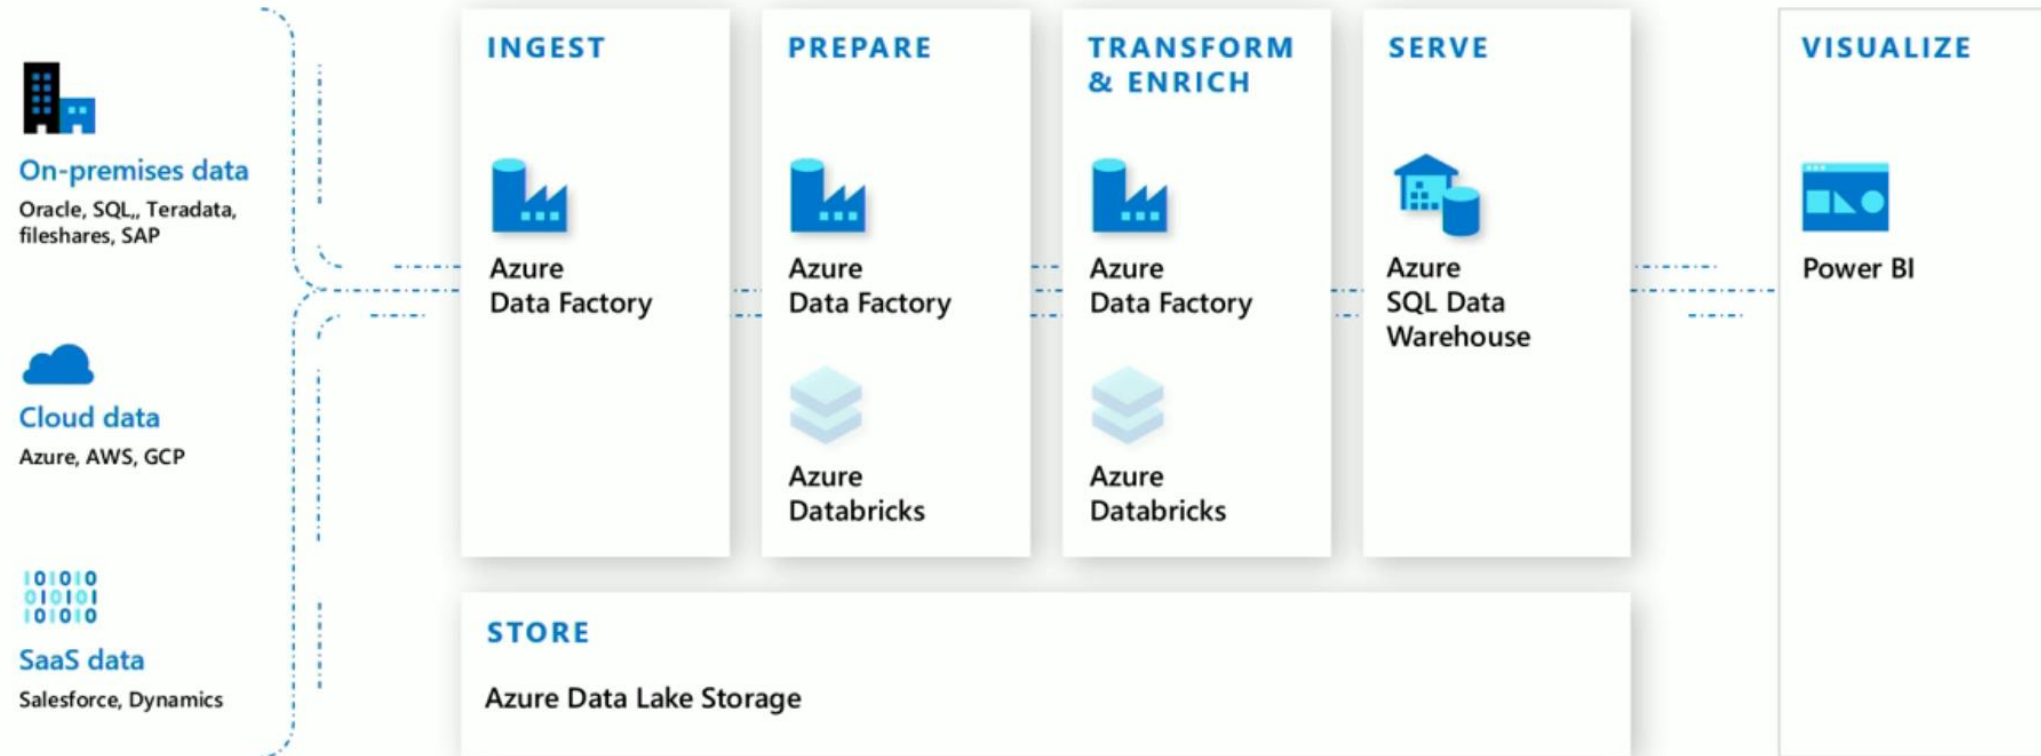


AGENDA




- A bit of theory & history...
- Let's practice: DEMO time
 - Delta Lake Tables in Lake Databases (Synapse)

Modern Data Warehouse



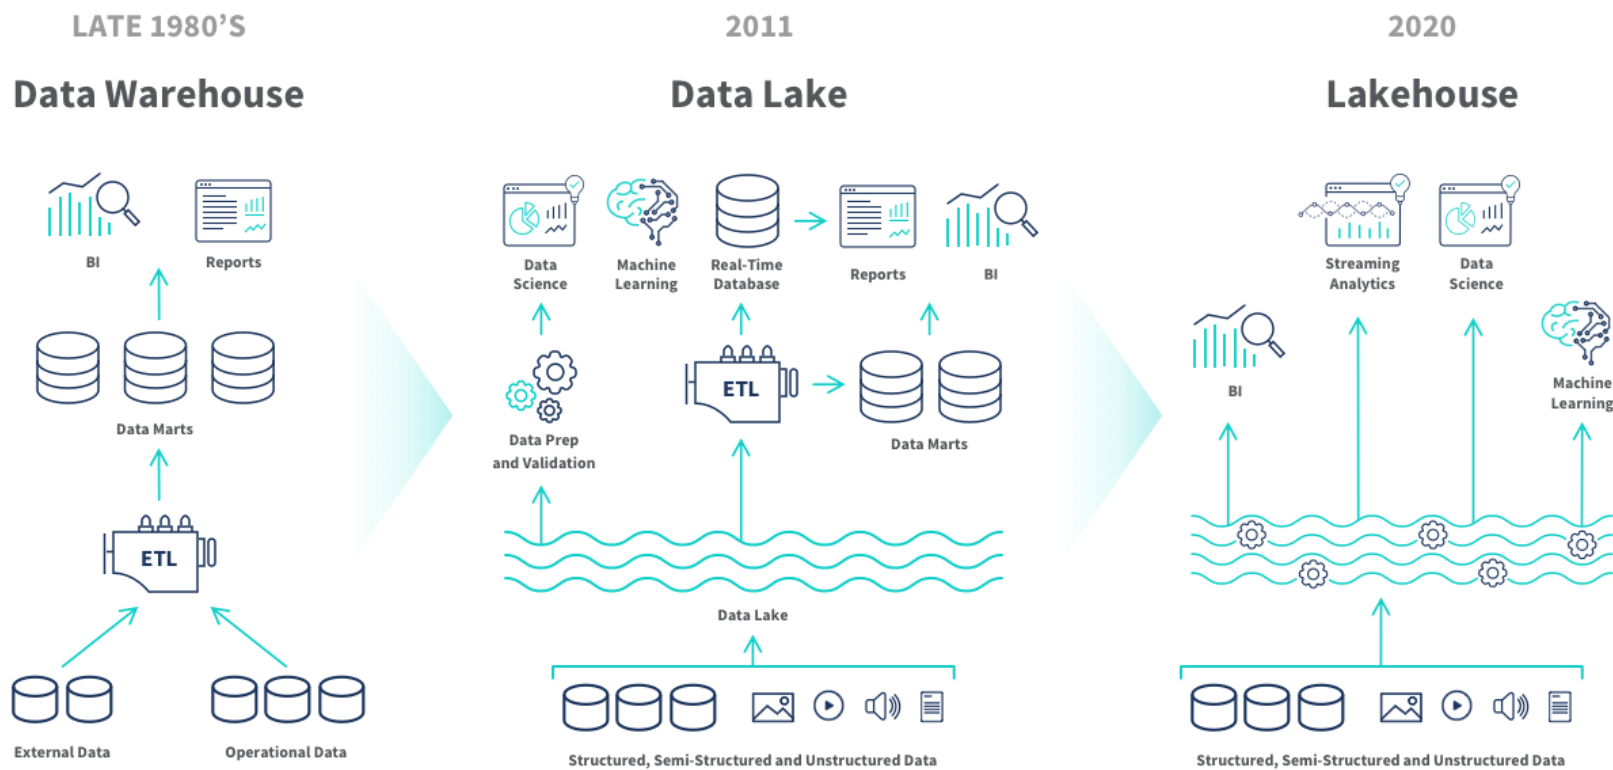
Data Lake + Warehouse = Lakehouse



> [What Is a Lakehouse?](#)



Data Warehouse evolution





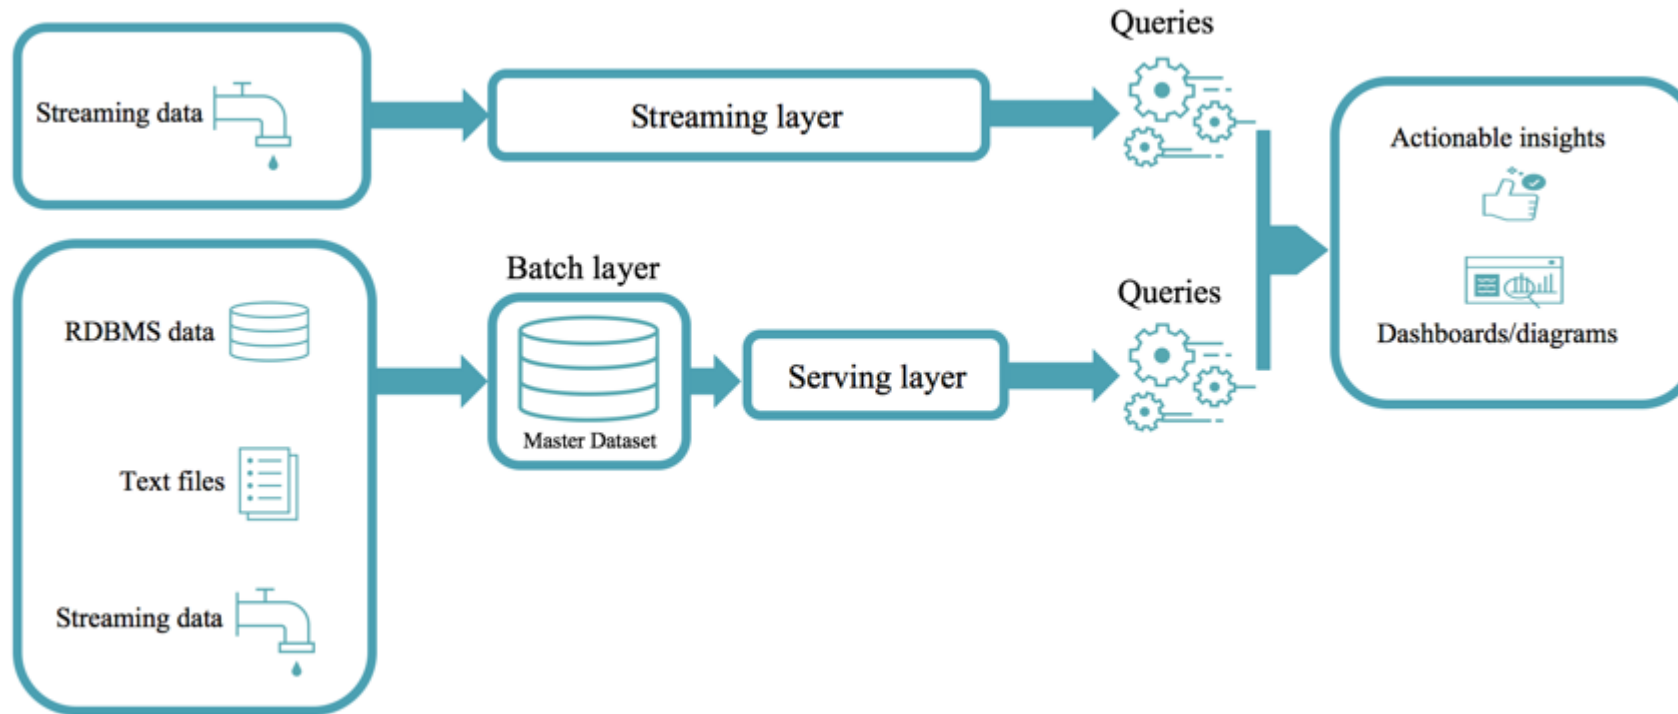
Introduction



- Lambda Architecture
- Streaming & batch
- Time travel
- Upserts & deletes



Lambda Architecture





Parquet format



What is Parquet?

- Apache Parquet is an **open source, column-oriented** data file format designed for efficient data storage and retrieval. It provides efficient data compression and encoding schemes with enhanced performance to handle complex data in bulk. Apache Parquet is designed to be a common interchange format for both batch and interactive workloads.

Characteristics of Parquet:

- Free and open source file format.
- Language agnostic.
- Column-based format
- Used for analytics (OLAP) use cases
- Highly efficient data compression and decompression.
- Supports complex data types and advanced nested data structures.

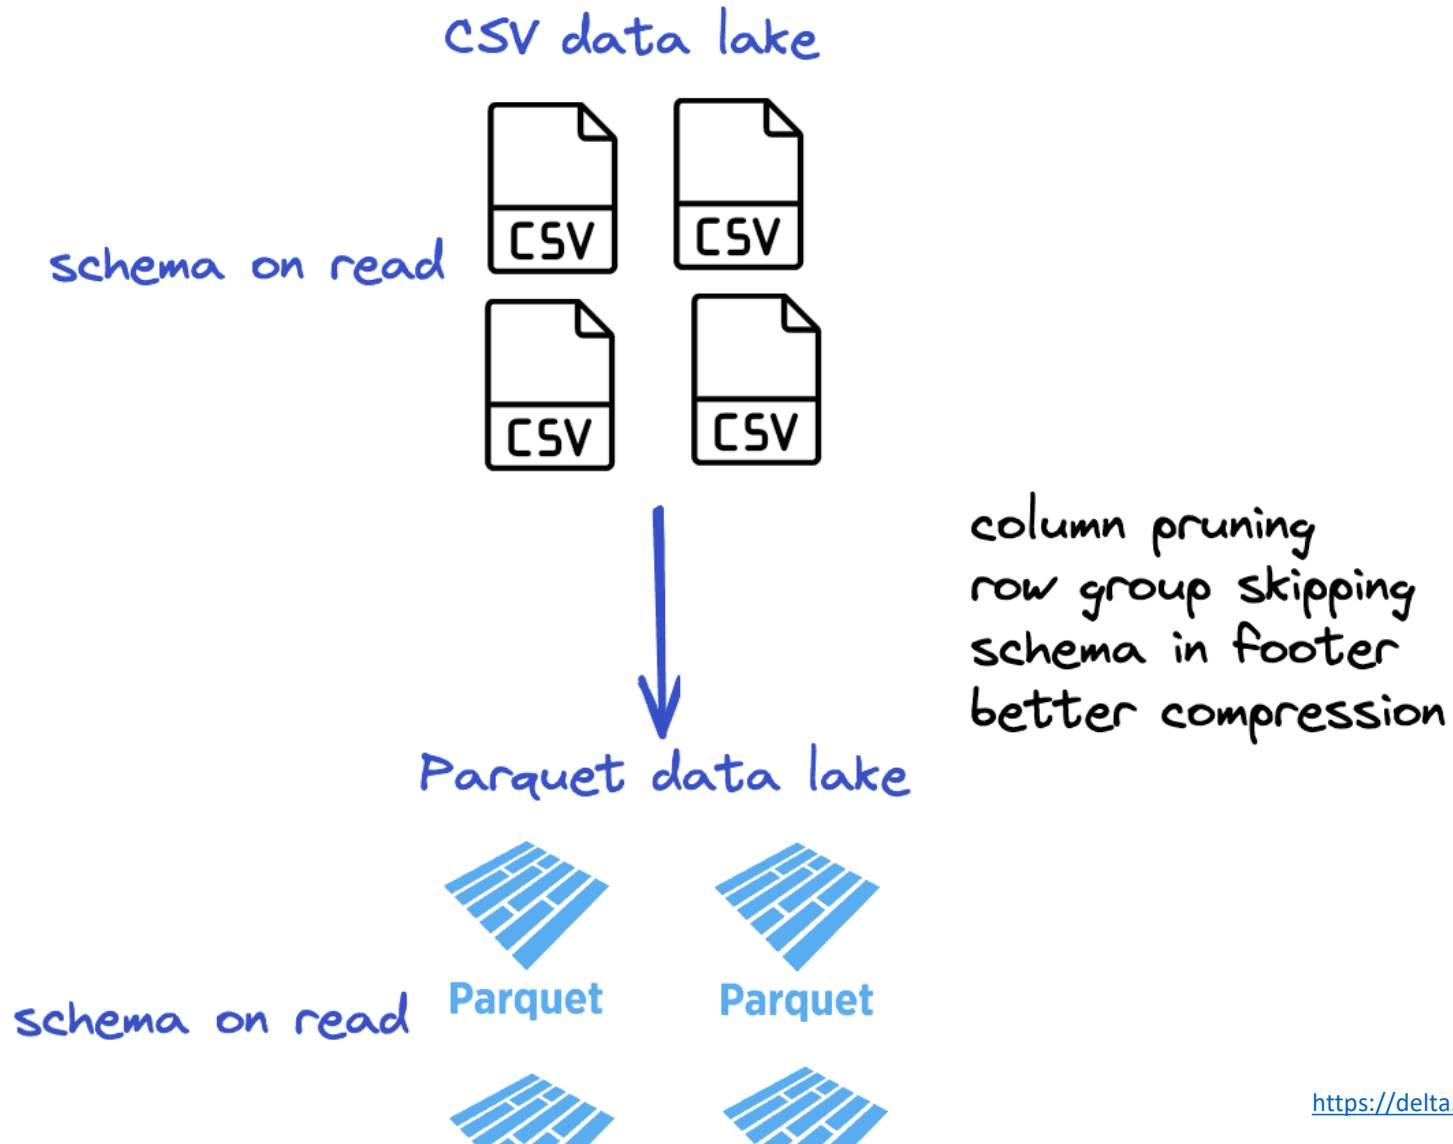


Problems with Parquet format



- No schema enforcement
- Updates – rewrite entire file
- Transactions inconsistency
- No interoperability between batch & streaming workloads
- No versioning

Advantages of Delta Lake over CSV





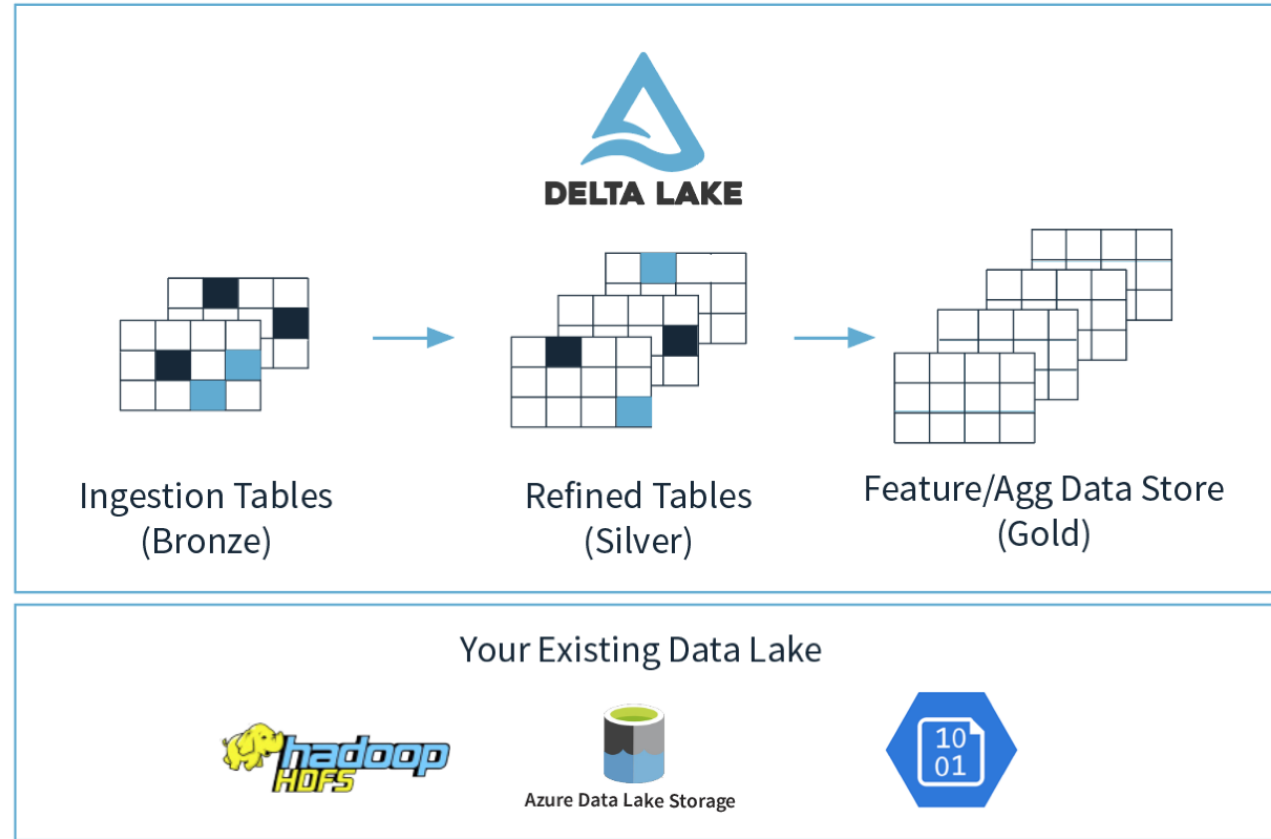
Medal architecture



Streaming



Batch



→ Analytics
and Machine
Learning



Delta – Key Features



ACID Transactions




Scalable
Metadata




Time Travel



Open Source




Unified
Batch/Streaming



Schema Evolution
/ Enforcement



Audit History



DML Operations



DEMO



Delta Lake Tables support



- Databricks (of course!)
- Delta format in Azure Data Factory & Synapse
 - Mapping Data Flow
 - Copy Activity (only with Databricks connector)
- Synapse notebook
- Synapse SQL Serverless (read – tables, views)
- Power BI



Schema Evolution

Delta Lake allows for schema evolution

1: Suppose you have this Delta table

first_name	age
bob	47
li	23
leonard	51



2: Data to append

first_name	age	country
frank	68	usa
jordana	26	brasil

→ Extra column of data

3: Schema enforcement will prevent the mismatched append

```
df.write.format("delta").mode("append").save("tmp/fun_people")
```

→ 🚔🛑 Exception: no data appended!

4: Use mergeSchema to enable schema evolution 🤗

```
df.write.option("mergeSchema", "true").mode("append").format("delta").save("tmp/fun_people")
```

```
spark.read.format("delta").load("tmp/fun_people").show()
```

first_name	age	country
jordana	26	brasil
frank	68	usa
leonard	51	null
bob	47	null
li	23	null

→ So amazingly cool!



DELTA LAKE

Press **Esc** to exit full screen

WHAT'S NEW? - RELEASE 2.3.0

`pip install delta-spark`



TABLE CLONE COMMANDS

```
-- Clone table definition (metadata)
CREATE TABLE taxi_clone LIKE taxi;

-- Clone table definition & data
CREATE TABLE taxi_clone
  SHALLOW CLONE taxi;
```

RECORD VACUUM OPERATIONS

```
-- Now recorded in the table history
VACUUM default.taxi;
DESCRIBE HISTORY default.taxi;
```

MERGE INTO COMMAND

```
/**
  Support when not
  matched by source
  clause
*/
MERGE INTO target AS t
  USING source AS s
    ON t.id = s.id
  WHEN MATCHED THEN
    UPDATE SET *
  WHEN NOT MATCHED BY
    SOURCE THEN DELETE;
```

READ TABLE CHANGES (CDF)

```
-- View changes between snapshot interval
SELECT * FROM table_changes('taxi', 3, 7);

-- Not using a Catalog? No problem.
SELECT * FROM table_changes_by_path(
  '/tmp/taxi', '2023-04-09', '2023-04-10');
```

CONVERT ICEBERG TABLES

```
-- Will not copy data (metadata op)
CONVERT TO DELTA
iceberg.`/tmp/iceberg/taxi`;
```



#1 - Zero-copy CONVERT TO DELTA



```
/**
```

```
Generates a Delta table in the same location.  
Will not copy data (metadata operation).
```

```
*/
```

```
CONVERT TO DELTA iceberg.`/tmp/iceberg/taxi`;
```




#2 - SHALLOW CLONE Command



```
/**
```

```
Clone a Delta, Parquet, or Iceberg table.  
Copies the table definition and data files.  
Note: creates a *new* transaction log.
```

```
*/
```

```
CREATE TABLE taxi_clone SHALLOW CLONE taxi;
```



#3 - Idempotent DML Operations



```
/**  
  Support idempotent writes for *all* DML operations  
  (INSERT/DELETE/UPDATE/MERGE).  
  `txnAppId` - a unique string that identifies the app  
  `txnVersion` - a monotonically increasing number  
*/  
SET spark.databricks.delta.write.txnAppId='DeltaLakeDemo';  
SET spark.databricks.delta.write.txnVersion=9446;  
UPDATE taxi  
  SET passenger_count=4  
  WHERE pride_id=1098;
```




#4 - WHEN NOT MATCHED BY SOURCE



```
/**  
    SQL support coming in Spark 3.4 (17 days)!  
*/
```

```
MERGE INTO default.taxi AS target  
USING new_taxi_data AS source  
    ON source.ride_id=target.ride_id  
WHEN NOT MATCHED BY SOURCE THEN DELETE;
```

```
# Use Python API in the meantime
```

```
targetDF.merge(sourceDF, "source.ride_id=target.ride_id") \  
    .whenNotMatchedBySourceDelete().execute()
```



#5 - CREATE TABLE LIKE Command



```
/**
```

```
Clone an existing Delta table.
```

```
Note:Copies the metadata only and no data files.
```

```
Will create an empty table.
```

```
*/
```

```
CREATE TABLE taxi_clone LIKE taxi;
```



#6 - Read Table Changes (CDF)



```
-- Read table changes that between an interval.
```

```
SELECT * FROM table_changes('taxi_clone', 3, 5);
```

```
-- Timestamps work too! (If you don't prefer versions)
```

```
SELECT * FROM table_changes('taxi_clone', '2023-04-09', '2023-04-10');
```

```
-- Don't have a Catalog? No problem!
```

```
SELECT * FROM table_changes_by_path('tmp/taxi_clone', '2023-04-09');
```



#7 - CDF for Column-mapped Tables



```
/**
  Previously, tables with column-mapping did *not* support
  change data feed reads after this operation.
*/
ALTER TABLE taxi RENAME COLUMN ride_id TO id; -- version 3

/**
  This release supports batch CDF reads for tables
  that have used column-mapping to DROP or RENAME a column.
*/
SELECT * FROM table_changes('taxi', 1, 3) -- this works now!
```




#8 - Faster S3 Reads & Writes



```
// Improves file listing efficiency during snapshot calc.  
// Note: this features works on S3A filesystems *only*.  
bin/spark-shell \  
  --packages io.delta:delta-core_2.12:2.3.0, \  
             org.apache.hadoop:hadoop-aws:3.3.1 \  
  --conf "spark.hadoop.delta.enableFastS3AListFrom=true"  
  
# Can be set the configuration on the Spark context too  
sc.hadoopConfiguration.set(  
  'spark.hadoop.delta.enableFastS3AListFrom', 'true')
```



#9 - Record VACUUM operations



```
/**  
    VACUUM operations and file metrics will now be recorded  
    in a table's history.  
*/  
VACUUM default.taxi;  
  
-- Start time, end time, and operation metrics will now appear!  
DESCRIBE HISTORY default.taxi;
```



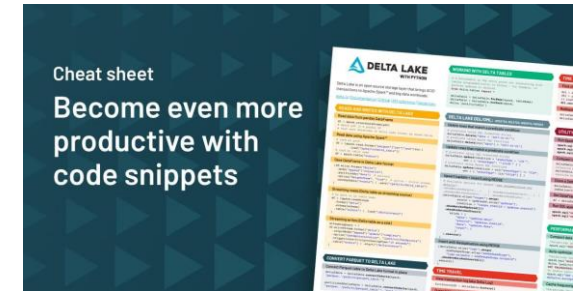
Summary



Resources



- <https://delta.io/>
- <https://docs.delta.io/latest/releases.html>
- <https://github.com/delta-io/delta/>
- <https://github.com/delta-io/delta/releases/>
- <https://github.com/delta-io/delta-examples/>
- [Table batch reads and writes](#)
- [Lambda Architecture](#)
- [Delta Lake Blogs](#)
- [What is Delta Lake? \(Microsoft\)](#)
- [Getting Started with Delta Lake \(Databricks\)](#)





Resources



DELTA LAKE 2.3.0

Support reading Change Data Feed (CDF) in SQL queries

Improved read and write performance on S3

Record **VACUUM** operations in the transaction log

Support reading Delta tables with deletion vectors



DELTA LAKE

Release Notes:

<https://github.com/delta-io/delta/releases/tag/v2.3.0>



@deltalakeoss



go.delta.io/slack



delta.io



Thank you!



kamil@azureplayer.net



@NowinskiK

@Azure_Player



AzurePlayer.net



<https://AzurePlayer.net/slides>

Kamil Nowinski

Microsoft Data Platform MVP

Analytics Architect, Azure DevOps Engineer Expert



15 edycja konferencji SQLDay

8-10 maja 2023, WROCŁAW + ONLINE



partner złoty

Future Processing

partner srebrny



partner brązowy

